

Open Resilient Cluster Manager (ORCM)

Ralph H. Castain, Ph.D.

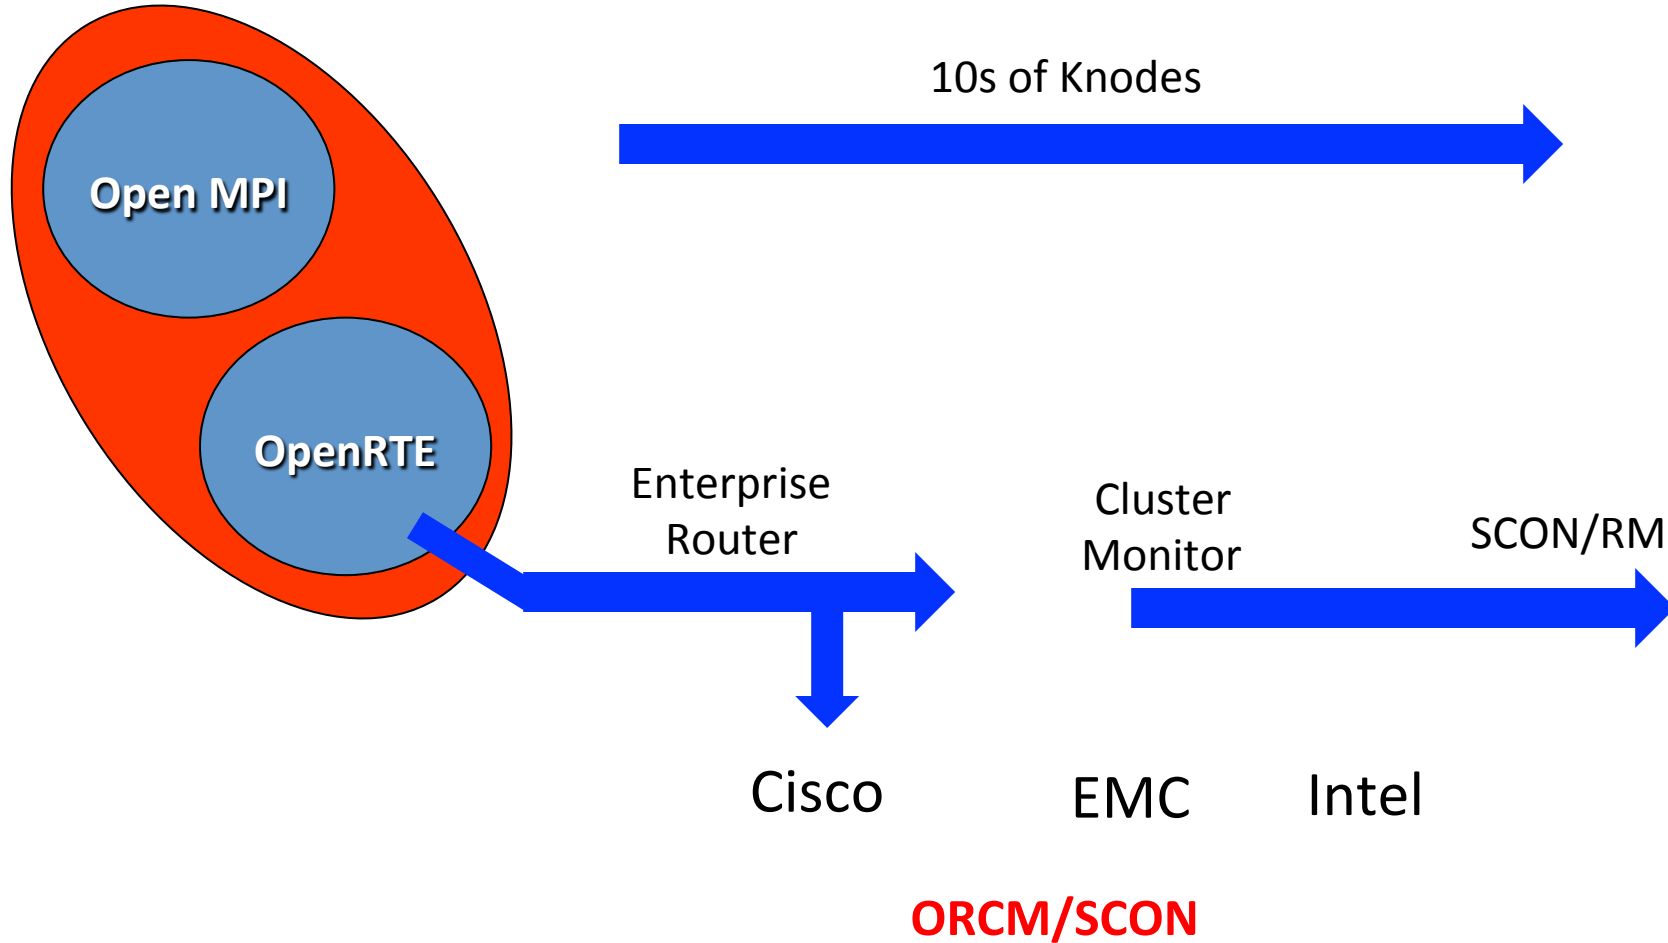
Objective

- Create a software platform
 - Easily customized, extended
 - Replace/override any behavior
 - Support proprietary as well as open extensions
 - Fully utilize existing installed infrastructure
- Establish ecosystem
 - Academic, industry collaborators, OMPI-like community
- Provide a reference solution
 - Publicly available, performant, flexible, scalable
 - Easily replace any pieces

2003 - present

OMPI/ORTE

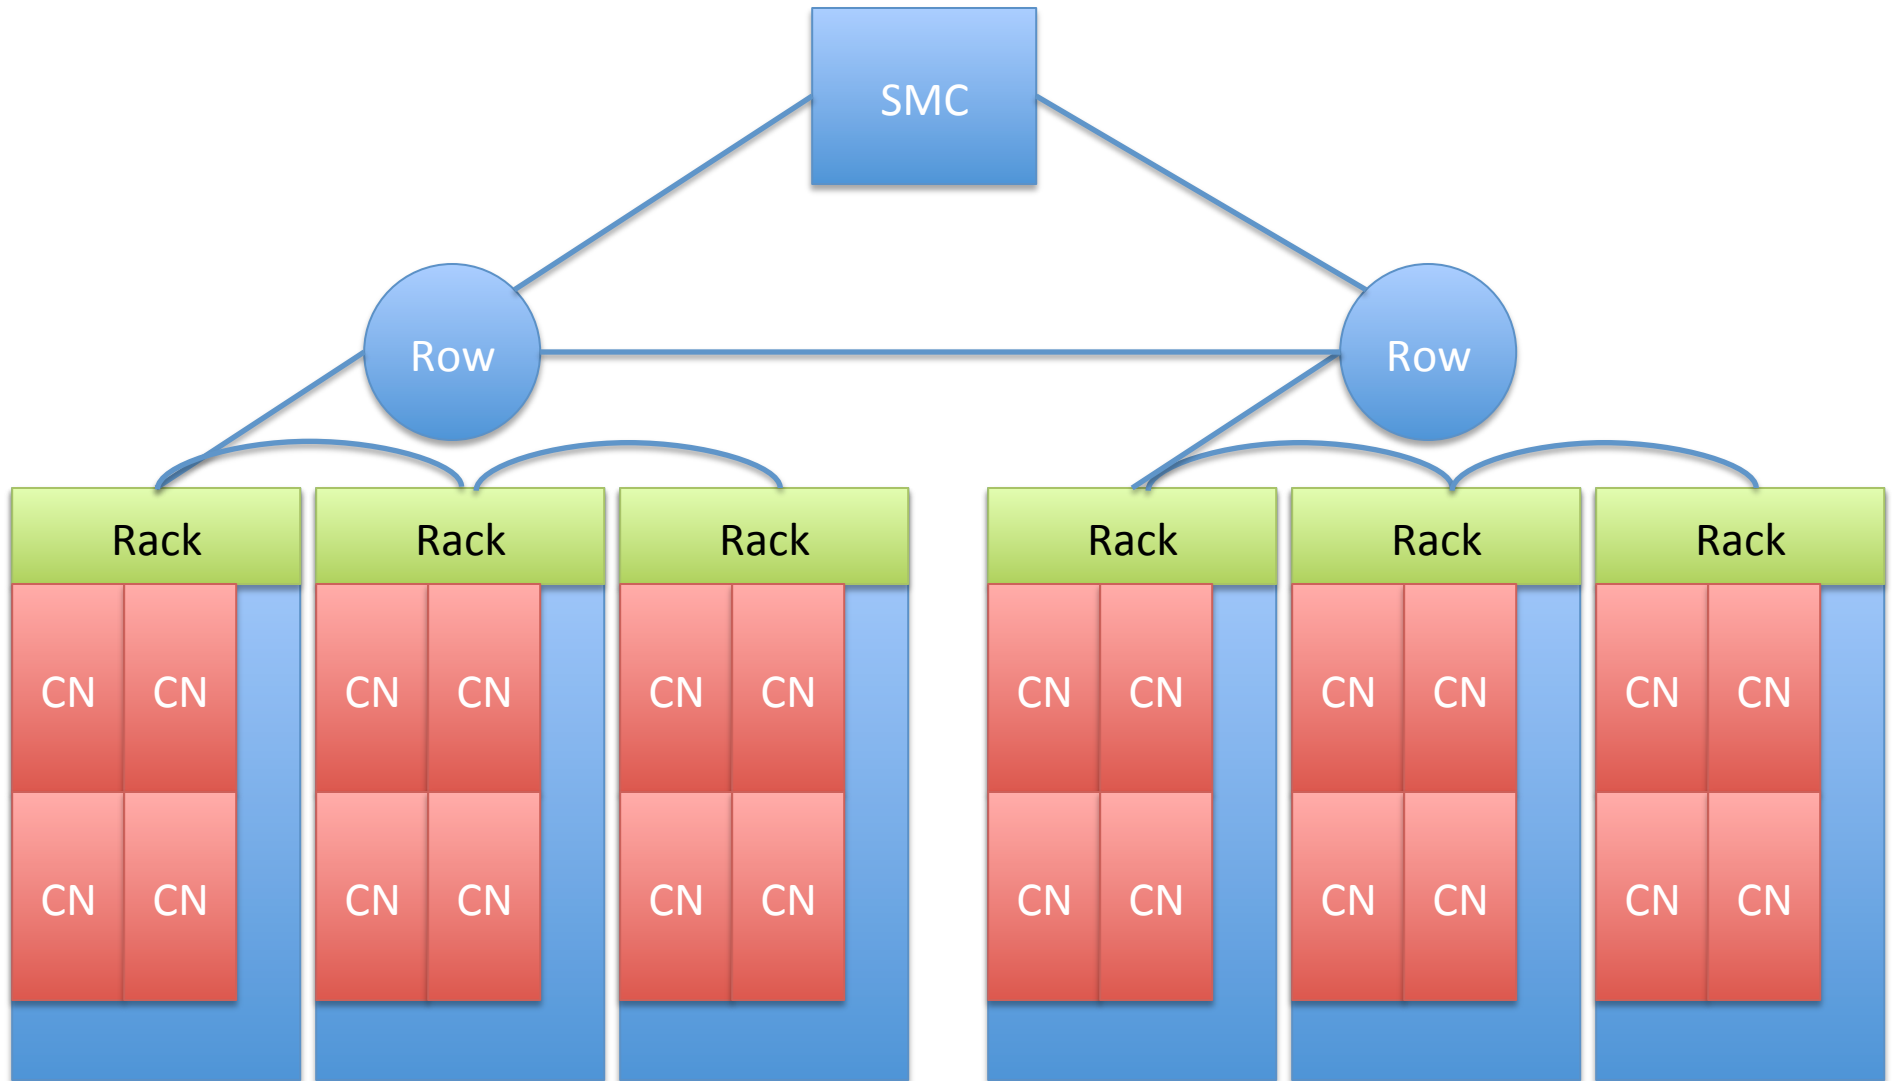
10s of Knodes



ORCM Roadmap

- Monitoring system (1Q2015)
 - System environment, power, process usage, etc.
- Overlay network/pub-sub (3Q2015)
 - More in a minute
- Job launch (3Q2015)
 - Can do it now, but need support for all MPIs
- Workload manager (2016)
 - Lightweight

Hierarchical Arch



Integration (examples)

- IO subsystem
 - Pre-stage executables and data
 - On-the-fly application-directed data positioning
 - Manage inter-node memory allocations, persistent NVM
- Network
 - QoS controls
 - Static endpoint support
- Power
 - Various modes, dynamic controls, site-level control

Instant On Steps

- Prestage executables to IO nodes
- Allocate and launch
 - Launch message => `orte_job_t`, included in allocation cmd
 - User-level step daemons spawned and initiate local launch
- Distributed mapping/rolling start (branch)
 - Each daemon computes map, stores all data (map, endpoints, network topo) in shared memory region for job
 - Connect/accept => pass SM connection
- Eliminate modex
 - Static endpoints provided by RM
 - (non-blocking) Lookup on first message
- Eliminate fence at end of `mpi_init`
 - Modex-recv becomes flag that proc is ready
 - RM flags all procs on node upon first request so subsequent checks are local

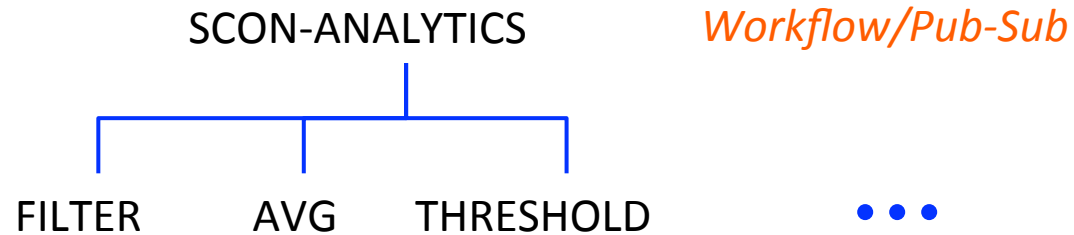
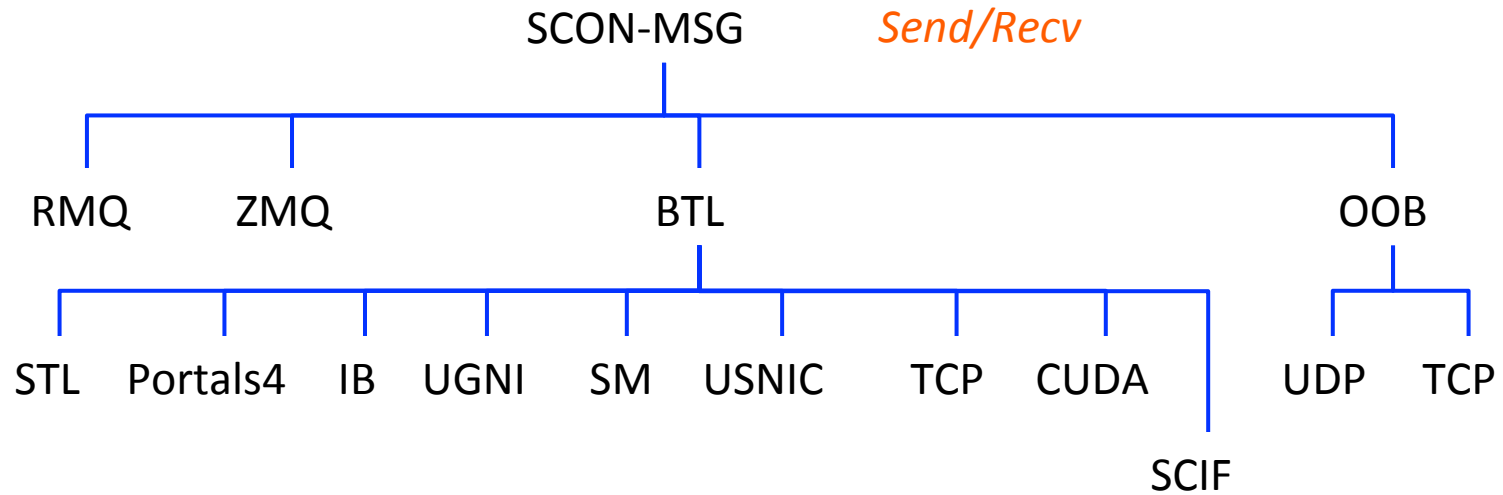
Definition: Overlay Network

- Messaging system
 - Scalable/resilient communications
 - Integration-friendly with user applications, system management software
- In-flight analytics
 - Insert analytic workflows anywhere in the data stream
 - Tap data stream at any point

Requirements

- Scalable to exascale levels
 - Better-than-linear scaling of broadcast
- Resilient
 - Self-heal around failures
 - Reintegrate recovered resources
- Dynamically configurable
 - Sense and adapt, user-directable
 - On-the-fly updates
- Open source (non-GPL)

High-Level Architecture



Messaging APIs

- Typical send/recv
 - Non-blocking, iovec or buffered (built-in heterogeneous support)
- Open channel
 - Specify remote peer and endpoint tag
 - Provide hints on type of data messaging to be used
 - Stream, command/control, etc.
 - Specify desired quality of service
 - Guaranteed delivery of every message, high priority, etc
- Subscribe to data stream
 - Specify source and data

Message Routing

- Detect/select
 - Various algorithms (trees, meshes)
 - Topological
- Defined per transport (branch)
- Heals routes
 - Provides alternate route upon failure
 - Up-level error if no alternate available on this transport
 - Allows re-routing over alternate transports (define prioritized policy)

Message Reliability

- Plugin architecture
 - Selected per transport, requested quality of service
- ACK-based (cmd/ctrl)
 - Ack each message, or window of messages, based on QoS
 - Resend or return error – QoS specified policy and number of retries before giving up
- NACK-based (streaming)
 - Nack if message sequence number is out of order indicating lost message(s)
 - Request resend or return error, based on QoS
- Multipath
 - Send each message across multiple paths
 - Recipient takes first received, discards rest

Analytics

- Event-based state machine
 - Each workflow in own thread, own instance of each plugin
 - Branch and merge of workflow
 - Tap stream between workflow steps
 - Tap data streams (sensors, others)
- Event generation
 - Generate events/alarms
 - Specify data to be included (window)